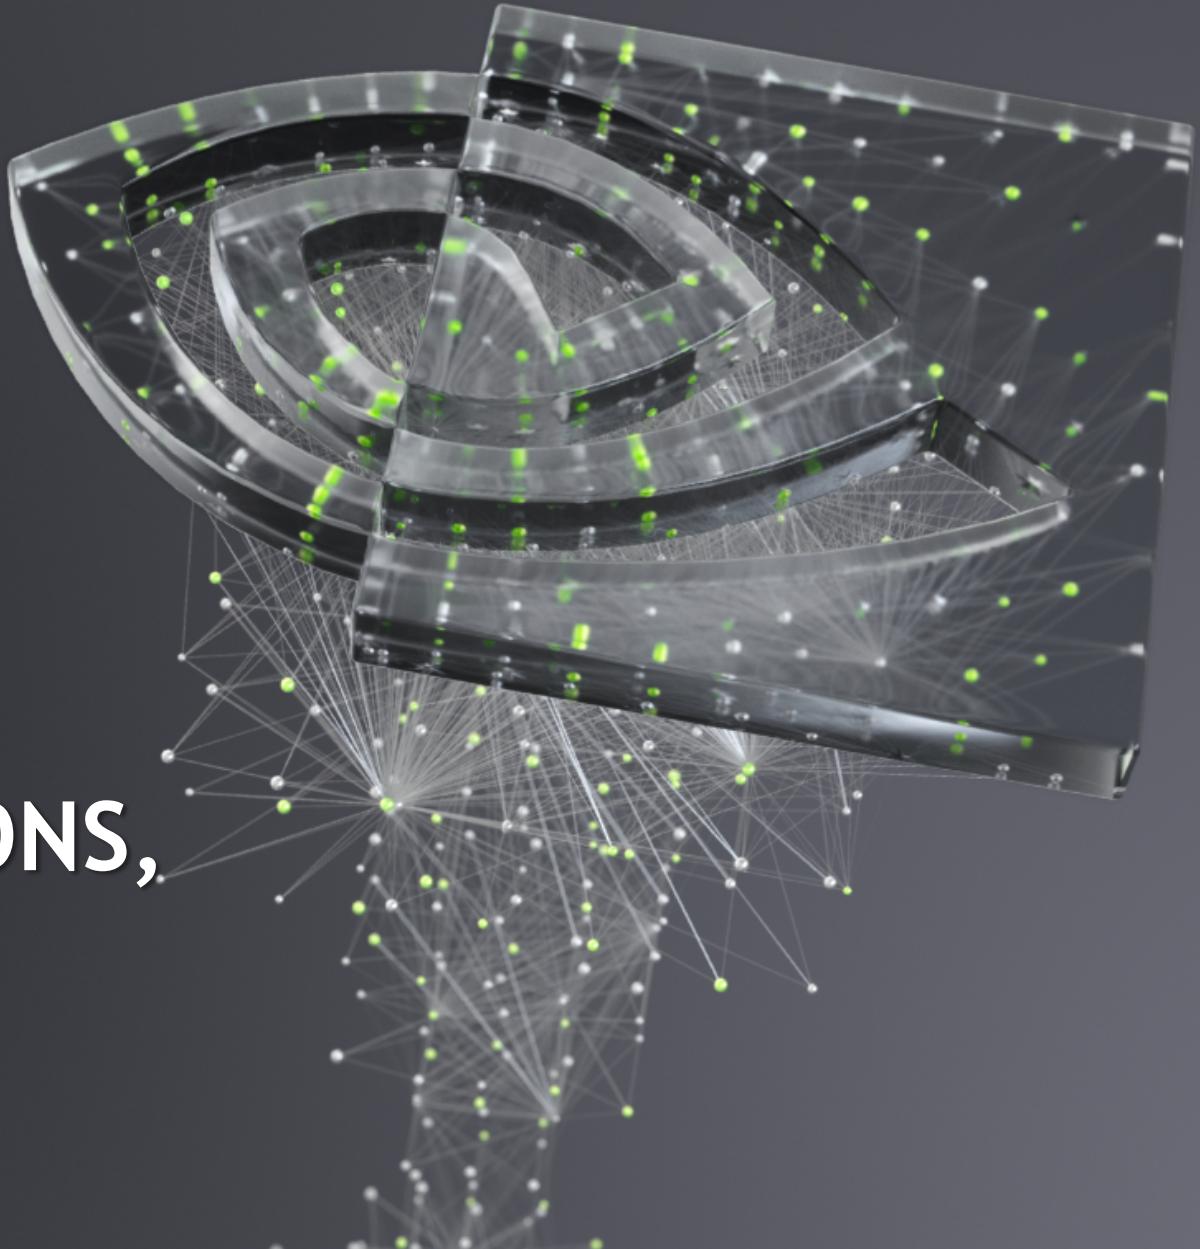




ATOMICS, REDUCTIONS, WARP SHUFFLE

Bob Crovella, 5/13/2020





AGENDA

- Transformations vs. Reductions, Thread Strategy
- Atomics, Atomic Reductions
- Atomic Tips and Tricks
- Classical Parallel Reduction
- Parallel Reduction + Atomics
- Warp Shuffle, Reduction with Warp Shuffle
- Other Warp Shuffle Uses
- Further Study
- Homework

A complex network graph is displayed against a dark gray background. The graph consists of numerous small, semi-transparent white and light green circular nodes, connected by thin, light gray lines representing edges. The nodes are distributed across the frame, with a higher density in the upper left and lower right quadrants, suggesting clusters or communities within the network.

ATOMICS

MOTIVATING EXAMPLE

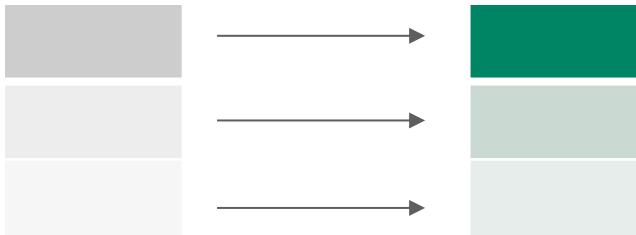
Sum - reduction

```
const int size = 100000;  
float a[size] = {...};  
float sum = 0;  
for (int i = 0; i < size; i++) sum += a[i];
```

-> sum variable contains the sum of all the elements of array a

TRANSFORMATION VS. REDUCTION

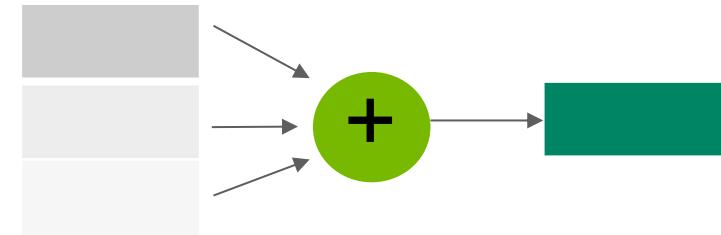
May guide the **thread strategy**: what will each thread do?



Transformation:

$$\text{e.g. } c[i] = a[i] + 10;$$

Thread strategy: one thread per output point



Reduction:

$$\text{e.g. } *c = \sum a[i]$$

Thread strategy: ??

REDUCTION: NAÏVE THREAD STRATEGY

One thread per input point

`*c += a[i];`

(Doesn't work.) Actual code the GPU executes:

LD R2, a[i] (thread independent)

LD R1, c (READ)

ADD R3, R1, R2 (MODIFY)

ST c, R3 (WRITE)

But **every thread** is trying to do this, potentially at the same time

The CUDA programming model does not enforce any order of thread execution

ATOMICS TO THE RESCUE

indivisible READ-MODIFY-WRITE

`atomicAdd(&c, a[i]);`

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>

LD R2, a[i]

(thread independent)

LD R1, c

(READ)

ADD R3, R1, R2

(MODIFY)

ST R3, c

(WRITE)

Becomes one indivisible operation/instruction:
RED.E.ADD.F32.FTZ.RN [c], R2;

Facilitated by special hardware in the L2 cache

May have performance implications

OTHER ATOMICS

- ▶ atomicMax/Min - choose the max (or min)
- ▶ atomicAdd/Sub - add to (or subtract from)
- ▶ atomicInc/Dec - increment (or decrement) and account for rollover/underflow
- ▶ atomicExch/CAS - swap values, or conditionally swap values
- ▶ atomicAnd/Or/Xor - bitwise ops
- ▶ atomics have different datatypes they can work on (e.g. int, unsigned, float, etc.)
- ▶ <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>

ATOMIC TIPS AND TRICKS

Determine my place in an order

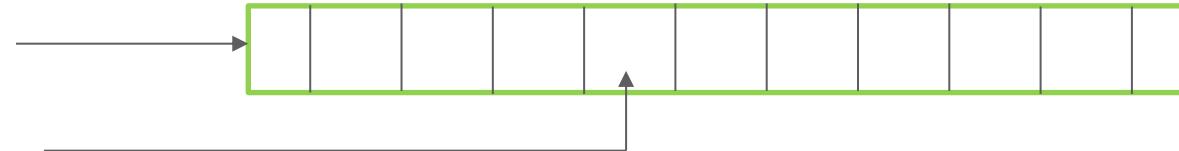
- ▶ Could be used to determine next work item, queue slot, etc.
- ▶ `int my_position = atomicAdd(order, 1);`
- ▶ Most atomics return a value that is the “old” value that was in the location receiving the atomic update.

ATOMIC TIPS AND TRICKS

Reserve space in a buffer

- ▶ Each thread in my kernel may produce a variable amount of data. How to collect all of this in one buffer, in parallel?

- ▶ `buffer_ptr:`



- ▶ `buffer_idx`

- ▶ `int my_dsize = var;`

- ▶ `float local_buffer[my_dsize] = {...};`

- ▶ `int my_offset = atomicAdd(buffer_idx, my_dsize);`

- ▶ `// buffer_ptr+my_offset now points to the first reserved location, of length my_dsize`

- ▶ `memcpy(buffer_ptr+my_offset, local_buffer, my_dsize*sizeof(float));`



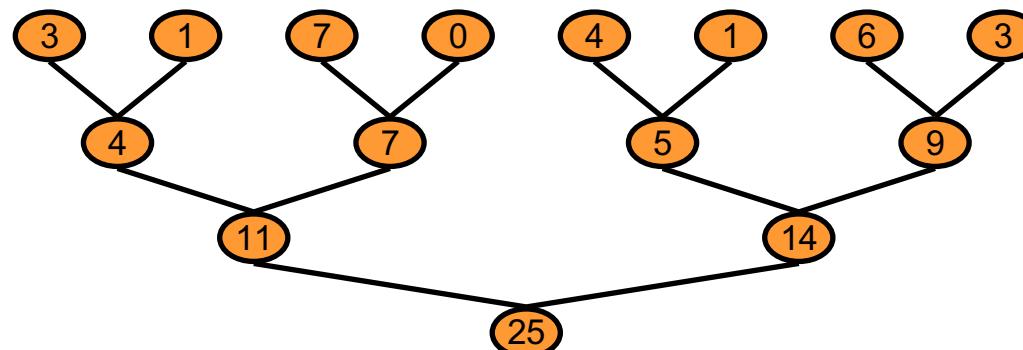
A network graph is displayed against a dark gray background. The graph consists of numerous small, semi-transparent white and light green circular nodes, connected by thin, light gray lines representing edges. The nodes are distributed across the frame, with a higher density in the upper half and some isolated or smaller clusters in the lower half. The overall effect is a complex web of connections, symbolizing a parallel computation or data distribution.

CLASSICAL PARALLEL REDUCTION

THE CLASSICAL PARALLEL REDUCTION

Atomics don't run at full memory bandwidth...

- ▶ We would like a reduction method that is not limited by atomic throughput
- ▶ We would like to effectively use all threads, as much as possible
- ▶ Parallel reduction is a common and important data parallel primitive
- ▶ Naïve implementations will often run into bottlenecks
- ▶ Basic methodology is a tree-based approach:



PROBLEM: GLOBAL SYNCHRONIZATION

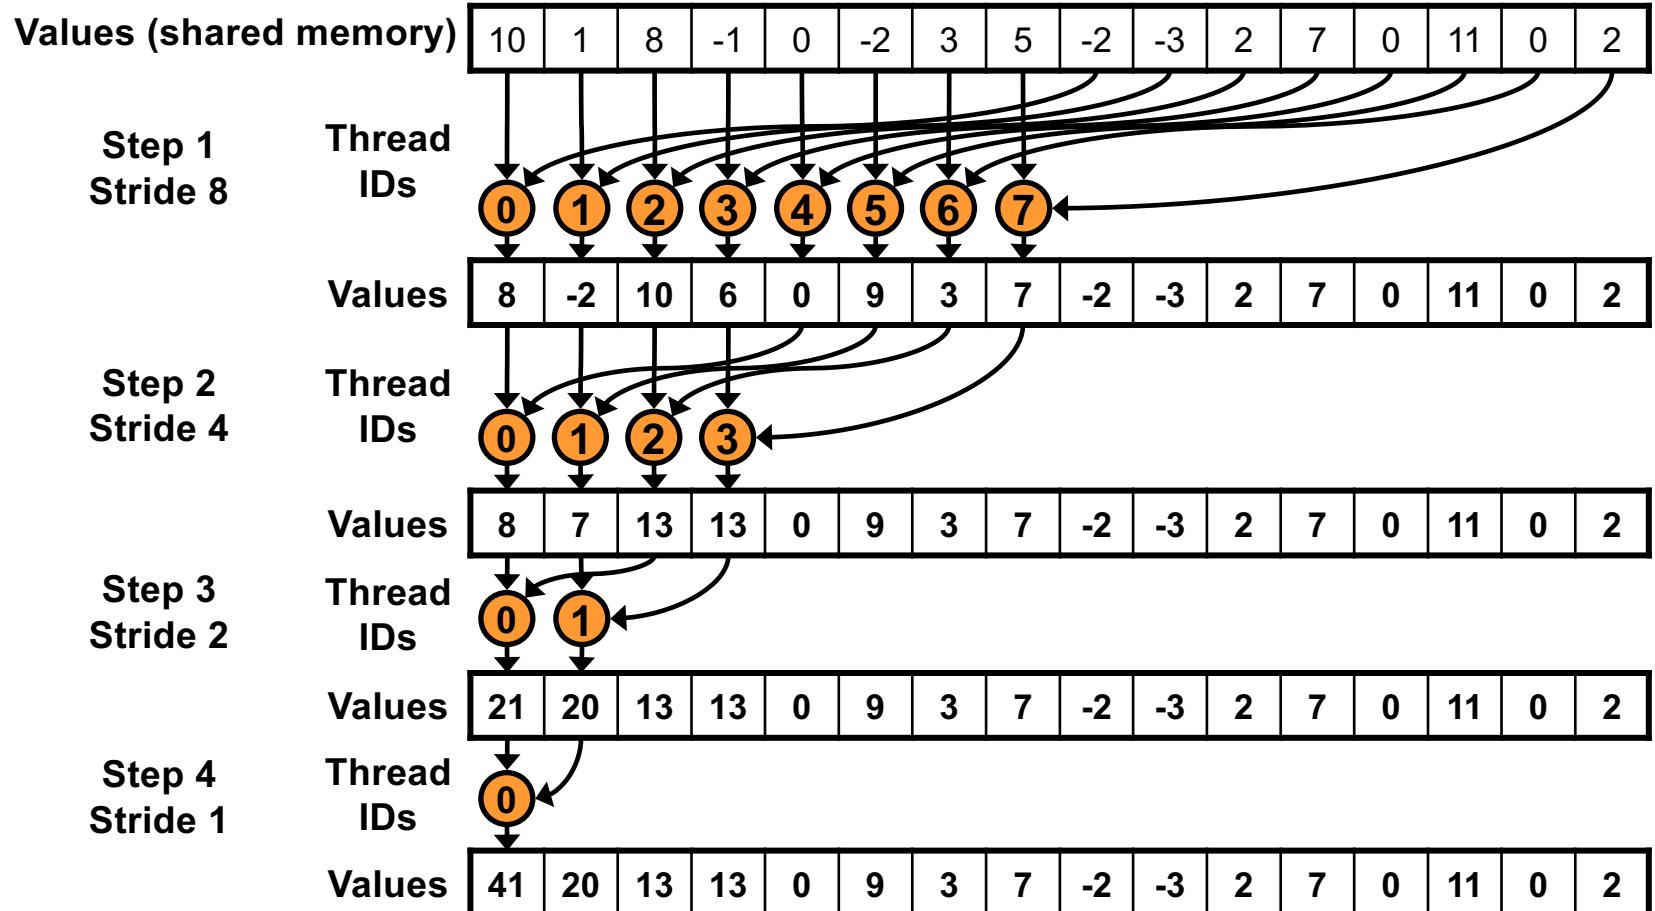
- ▶ If we could synchronize across all thread blocks, could easily reduce very large arrays, right?
 - ▶ Global sync after each block produces its result
 - ▶ Once all blocks reach sync, continue recursively
- ▶ One possible solution: decompose into multiple kernels
 - ▶ Kernel launch serves as a global synchronization point
 - ▶ Kernel launch has low SW overhead (but not zero)
- ▶ Other possible solutions:
 - ▶ Use atomics at the end of threadblock-level reduction
 - ▶ Use a threadblock-draining approach (see `threadFenceReduction` sample code)
 - ▶ Use cooperative groups - cooperative kernel launch

```

for (unsigned int s=blockDim.x/2; s>0; s>>=1) {
    if (tid < s) {
        sdata[tid] += sdata[tid + s];
    }
    _syncthreads(); // outside the if-statement
}

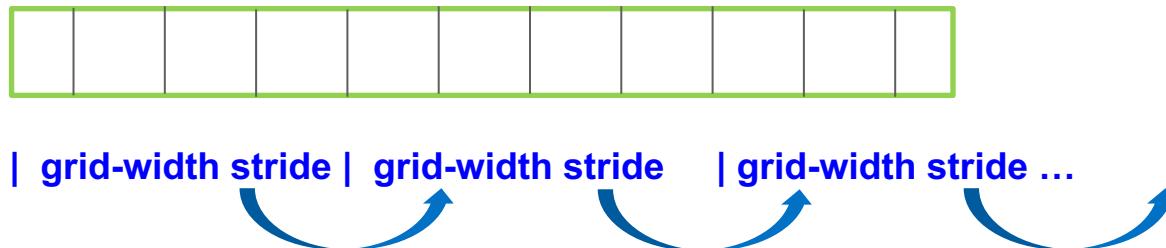
```

SEQUENTIAL ADDRESSING



DETOUR: GRID-STRIDE LOOPS

- ▶ We'd like to be able to design kernels that load and operate on arbitrary data sizes efficiently
- ▶ Want to maintain coalesced loads/stores, efficient use of shared memory
- ▶ Can also be used for ninja-level tuning - choose number of blocks sized to the GPU
- ▶ `gdata[0..N-1]`:



```
int idx = threadIdx.x + blockDim.x * blockIdx.x;
while (idx < N) {
    sdata[tid] += gdata[idx];
    idx += gridDim.x * blockDim.x; // grid width
}
```

PUTTING IT ALL TOGETHER

```
__global__ void reduce(float *gdata, float *out){  
    __shared__ float sdata[BLOCK_SIZE];  
    int tid = threadIdx.x;  
    sdata[tid] = 0.0f;  
    size_t idx = threadIdx.x+blockDim.x*blockIdx.x;  
  
    while (idx < N) { // grid stride loop to load data  
        sdata[tid] += gdata[idx];  
        idx += blockDim.x*blockDim.x;  
    }  
  
    for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
        __syncthreads();  
        if (tid < s) // parallel sweep reduction  
            sdata[tid] += sdata[tid + s];  
    }  
    if (tid == 0) out[blockIdx.x] = sdata[0];  
}
```

GETTING RID OF THE 2ND KERNEL CALL

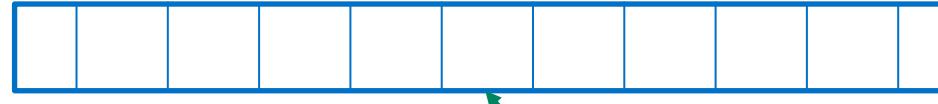
```
__global__ void reduce_a(float *gdata, float *out){  
    __shared__ float sdata[BLOCK_SIZE];  
    int tid = threadIdx.x;  
    sdata[tid] = 0.0f;  
    size_t idx = threadIdx.x+blockDim.x*blockIdx.x;  
  
    while (idx < N) { // grid stride loop to load data  
        sdata[tid] += gdata[idx];  
        idx += blockDim.x*blockDim.x;  
    }  
  
    for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
        __syncthreads();  
        if (tid < s) // parallel sweep reduction  
            sdata[tid] += sdata[tid + s];  
    }  
    if (tid == 0) atomicAdd(out, sdata[0]);  
}
```

The background of the image features a complex network graph. It consists of numerous small, semi-transparent white and light green circular nodes scattered across a dark gray background. These nodes are interconnected by a dense web of thin, light gray lines, representing connections or relationships between the data points.

WARP SHUFFLE

INTER-THREAD COMMUNICATION: SO FAR

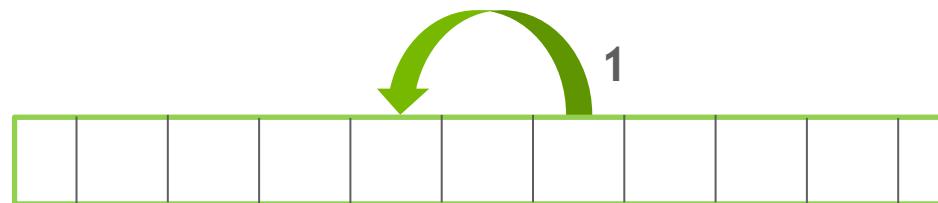
- ▶ Using shared memory:



- ▶ Threads:



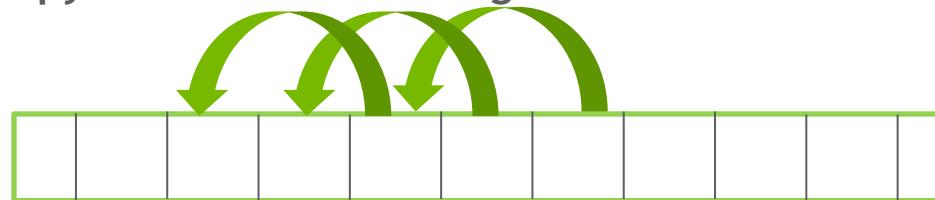
- ▶ Wouldn't this be convenient:



- ▶ Threads:

INTRODUCING WARP SHUFFLE

- ▶ Allows for intra-warp communication
- ▶ Various supported movement patterns:
 - ▶ `__shfl_sync()`: copy from lane ID (arbitrary pattern)
 - ▶ `__shfl_xor_sync()`: copy from calculated lane ID (calculated pattern)
 - ▶ `__shfl_up_sync()`: copy from delta/offset lower lane
 - ▶ `__shfl_down_sync()`: copy from delta/offset higher lane:



- ▶ Both source and destination threads in the warp must “participate”
- ▶ Sync “mask” used to identify and reconverge needed threads

WARP SHUFFLE REDUCTION

```
__global__ void reduce_ws(float *gdata, float *out){  
    __shared__ float sdata[32];  
    int tid = threadIdx.x;  
    int idx = threadIdx.x+blockDim.x*blockIdx.x;  
    float val = 0.0f;  
    unsigned mask = 0xFFFFFFFFU;  
    int lane = threadIdx.x % warpSize;  
    int warpID = threadIdx.x / warpSize;  
    while (idx < N) { // grid stride loop to load  
        val += gdata[idx];  
        idx += blockDim.x*blockDim.x;  
    }  
  
    // 1st warp-shuffle reduction  
    for (int offset = warpSize/2; offset > 0; offset >>= 1)  
        val += __shfl_down_sync(mask, val, offset);  
    if (lane == 0) sdata[warpID] = val;  
    __syncthreads(); // put warp results in shared mem
```

```
// hereafter, just warp 0  
    if (warpID == 0){  
        // reload val from shared mem if warp existed  
        val = (tid < blockDim.x/warpSize)?sdata[lane]:0;  
  
        // final warp-shuffle reduction  
        for (int offset = warpSize/2; offset > 0; offset >>= 1)  
            val += __shfl_down_sync(mask, val, offset);  
  
        if (tid == 0) atomicAdd(out, val);  
    }  
}
```

WARP SHUFFLE BENEFITS

- ▶ Reduce or eliminate shared memory usage
- ▶ Single instruction vs. 2 or more instructions
- ▶ Reduce level of explicit synchronization

WARP SHUFFLE TIPS AND TRICKS

What else can we do with it?

- ▶ Broadcast a value to all threads in the warp in a single instruction
- ▶ Perform a warp-level prefix sum
- ▶ Atomic aggregation

FUTURE SESSIONS

- ▶ Using Managed Memory
- ▶ Concurrency (streams, copy/compute overlap, multi-GPU)
- ▶ Analysis Driven Optimization
- ▶ Cooperative Groups

FURTHER STUDY

- ▶ Parallel reduction:
 - ▶ <https://developer.download.nvidia.com/assets/cuda/files/reduction.pdf>
- ▶ Warp-shuffle and reduction:
 - ▶ <https://devblogs.nvidia.com/faster-parallel-reductions-kepler/>
- ▶ CUDA Cooperative Groups:
 - ▶ <https://devblogs.nvidia.com/cooperative-groups/>
- ▶ Grid-stride loops:
 - ▶ <https://devblogs.nvidia.com/cuda-pro-tip-write-flexible-kernels-grid-stride-loops/>
- ▶ Floating point:
 - ▶ <https://developer.nvidia.com/sites/default/files/akamai/cuda/files/NVIDIA-CUDA-Floating-Point.pdf>
- ▶ CUDA Sample Codes:
 - ▶ Reduction, threadFenceReduction, reductionMultiBlockCG

HOMEWORK

- ▶ Log into Summit (ssh username@home.ccs.ornl.gov -> ssh summit)
- ▶ Clone GitHub repository:
 - ▶ Git clone [git@github.com:olcf/cuda-training-series.git](https://github.com/olcf/cuda-training-series.git)
- ▶ Follow the instructions in the readme.md file:
 - ▶ <https://github.com/olcf/cuda-training-series/blob/master/exercises/hw5/readme.md>
- ▶ Prerequisites: basic linux skills, e.g. ls, cd, etc., knowledge of a text editor like vi/emacs, and some knowledge of C/C++ programming

